

Statistical Inference: matched pairs and normal quantile plot

Matched pairs

Some data:

subject	druga	drugb
1	2.0	3.5
2	3.6	5.7
3	2.6	2.9
4	2.6	2.4
5	7.3	9.9
6	3.4	3.3
7	14.9	16.7
8	6.6	6.0
9	2.3	3.8
10	2.0	4.0
11	6.8	9.1
12	8.5	20.9

Matched pairs data

- Data are comparison of 2 drugs for effectiveness at reducing pain.
 - 12 subjects (cases) were arthritis sufferers
 - Response is #hours of pain relief from each drug.
- In reading example, each child tried only one reading method.
- But here, each subject tried out both drugs, giving us two measurements.
 - Possible because, if you wait long enough, one drug has no influence over effect of other.
 - Advantage: focused comparison of drugs. Compare one drug with another on same person, removes a lot of variability due to differences between people.
 - Matched pairs, requires different analysis.
- Design: randomly choose 6 of 12 subjects to get drug A first, other 6 get drug B first.

Paired t test: reading the data

Values aligned in columns:

```
my_url <- "http://www.utsc.utoronto.ca/~butler/c32/analgesic.t  
pain <- read_table(my_url)
```

```
## Parsed with column specification:  
## cols(  
##   subject = col_double(),  
##   druga = col_double(),  
##   drugb = col_double()  
## )
```

The data

pain

subject	druga	drugb
1	2.0	3.5
2	3.6	5.7
3	2.6	2.9
4	2.6	2.4
5	7.3	9.9
6	3.4	3.3
7	14.9	16.7
8	6.6	6.0
9	2.3	3.8
10	2.0	4.0
11	6.8	9.1
12	8.5	20.9

Paired t -test

```
with(pain, t.test(druga, drugb, paired = T))
```

```
##  
## Paired t-test  
##  
## data:  druga and drugb  
## t = -2.1677, df = 11, p-value = 0.05299  
## alternative hypothesis: true difference in means is not equal to  
## 95 percent confidence interval:  
## -4.29941513  0.03274847  
## sample estimates:  
## mean of the differences  
## -2.133333
```

- P-value is 0.053.
- Not quite evidence of difference between drugs.

t-testing the differences

- Likewise, you can calculate the differences yourself and do a 1-sample t-test on them.
- First calculate a column of differences:

```
(pain %>% mutate(diff=druga-drugb) -> pain)
```

subject	druga	drugb	diff
1	2.0	3.5	-1.5
2	3.6	5.7	-2.1
3	2.6	2.9	-0.3
4	2.6	2.4	0.2
5	7.3	9.9	-2.6
6	3.4	3.3	0.1
7	14.9	16.7	-1.8
8	6.6	6.0	0.6
9	2.3	3.8	-1.5
10	2.0	4.0	-2.0
11	6.8	9.1	-2.3
12	8.5	20.9	-12.4

t-test on the differences

- then throw them into `t.test`, testing that the mean is zero, with same result as before:

```
with(pain,t.test(diff,mu=0))
```

```
##  
## One Sample t-test  
##  
## data: diff  
## t = -2.1677, df = 11, p-value = 0.05299  
## alternative hypothesis: true mean is not equal to 0  
## 95 percent confidence interval:  
## -4.29941513 0.03274847  
## sample estimates:  
## mean of x  
## -2.133333
```



Assessing normality

- 1-sample and 2-sample t-tests assume (each) group normally distributed.
- Matched pairs analyses assume (theoretically) that differences normally distributed.
- Though we know that t-tests generally behave well even without normality.
- How to assess normality? A normal quantile plot.
 - Idea: scatter of points should follow the straight line, without curving.
 - Outliers show up at bottom left or top right of plot as points off the line.

The normal quantile plot

- of differences from matched pairs data

```
ggplot(pain, aes(sample=diff)) + stat_qq() + stat_qq_line()
```




- Points should follow the straight line. Bottom left one way off, so normality questionable here: outlier.

More normal quantile plots

- How straight does a normal quantile plot have to be?
- There is randomness in real data, so even a normal quantile plot from normal data won't look perfectly straight.
- With a small sample, can look not very straight even from normal data.
- Looking for systematic departure from a straight line; random wiggles ought not to concern us.
- Look at some examples where we know the answer, so that we can see what to expect.


Normal data, large sample

```
d=tibble(x=rnorm(200))  
ggplot(d,aes(x=x))+geom_histogram(bins=10)
```



The normal quantile plot


```
ggplot(d,aes(sample=x))+stat_qq()+stat_qq_line()
```



Normal data, small sample

- Not so convincingly normal, but not obviously skewed:


```
d=tibble(x=rnorm(20))  
ggplot(d,aes(x=x))+geom_histogram(bins=5)
```



The normal quantile plot

Good, apart from the highest and lowest points being slightly off. I'd call this good:


```
ggplot(d,aes(sample=x))+stat_qq()+stat_qq_line()
```



Chi-squared data, $df = 10$

Somewhat skewed to right:


```
d=tibble(x=rchisq(100,10))  
ggplot(d,aes(x=x))+geom_histogram(bins=10)
```



The normal quantile plot

Somewhat opening-up curve:


```
ggplot(d,aes(sample=x))+stat_qq()+stat_qq_line()
```



Chi-squared data, $df = 3$

Definitely skewed to right:


```
d=tibble(x=rchisq(100,3))  
ggplot(d,aes(x=x))+geom_histogram(bins=10)
```



The normal quantile plot

Clear upward-opening curve:


```
ggplot(d,aes(sample=x))+stat_qq()+stat_qq_line()
```



t-distributed data, $df = 3$

Long tails (or a very sharp peak):


```
d=tibble(x=rt(300,3))  
ggplot(d,aes(x=x))+geom_histogram(bins=10)
```



The normal quantile plot


Low values too low and high values too high for normal.

```
ggplot(d,aes(sample=x))+stat_qq()+stat_qq_line()
```



Our pain-relief data

```
ggplot(pain, aes(sample=diff)) + stat_qq() + stat_qq_line()
```



Comments

- Definitely not normal. What to do?
- Sign test on differences, null median 0.

Sign test

- Most easily: calculate differences in data frame, then use `smmr`.
- Null median difference is 0:

```
pain %>% mutate(mydiff=druga-drugb) %>%  
sign_test(mydiff,0)
```

```
## $above_below  
## below above  
##      9      3  
##  
## $p_values  
##   alternative    p_value  
## 1         lower 0.07299805  
## 2          upper 0.98071289  
## 3    two-sided 0.14599609
```


Comments

- P-value 0.1460. No evidence that the drugs are different.
- Since we are working in a pipeline, input data frame to `sign_test` is “whatever came out of previous step”.

(Some of) the kids' reading data, again

```
## Parsed with column specification:
## cols(
##   group = col_character(),
##   score = col_double()
## )

kids %>% slice_sample(n=12)
```


group	score
t	49
c	48
c	46
c	17
t	43
t	44
c	26
.	50

Where we were at

- 21 kids in “treatment”, new reading method; 23 in “control”, standard reading method.
- Assessing assumptions:
 - We did two-sample t-test (Satterthwaite-Welch) before.
 - Assumes approx. normal data within each group.
 - Does not assume equal spread.
 - (Pooled t-test *does* assume equal spread).
 - Assess each group separately.

Boxplots for reading data


```
ggplot(kids, aes(x=group, y=score)) + geom_boxplot()
```



Facetted normal quantile plots

Done this way:

```
ggplot(kids, aes(sample=score)) + stat_qq() + stat_qq_line() +  
facet_wrap(~group)
```



Comments

- These plots show no problems with normality. Both groups are more or less symmetric/normal and there are no outliers.
- Equal spreads questionable, but we don't need that.
- Assess equal spreads by looking at *slopes* of normal quantile plots.
- We ought be happy with the (Welch) two-sample t-test (over)

Welch two-sample test

```
t.test(score~group,data=kids,alternative="less")
```

```
##  
##  Welch Two Sample t-test  
##  
## data:  score by group  
## t = -2.3109, df = 37.855, p-value = 0.01319  
## alternative hypothesis: true difference in means is less than 0  
## 95 percent confidence interval:  
##      -Inf -2.691293  
## sample estimates:  
## mean in group c mean in group t  
##      41.52174      51.47619
```

from which we concluded that the new reading method really does help.